

# **LLM Project Report**

---

## **1. Student Information**

- Student Name:** Beneetha Thattayath Ambika
  - Student ID (optional):** S224154192
  - Date Submitted:** 30/05/2025
- 

## **2. Project Introduction**

- Title of the Project:** Fine-tuning LLM for Enterprise Applications
- What is the project about?**

This project focuses on training a large language model to detect and classify medical misinformation. The Mistral-7B-Instruct-v0.1 model was fine-tuned in two stages using datasets such as HealthFact, SciFact, and COVID-19 Fake News. The objective is to accurately categorize health-related claims into **TRUE**, **FALSE**, or **MISLEADING**.

Prior to fine-tuning, prompt formats and hyperparameters were optimized to improve model response quality. The two-stage approach allowed the model to learn both general medical knowledge and domain-specific misinformation patterns. A Gradio-based user interface was also developed, enabling users to input claims and receive predictions along with a trust score based on semantic similarity.

The project aims to improve the trustworthiness and effectiveness of AI systems in handling health-related information.

- Why is this project important or useful?**

This project is important because it addresses the growing problem of health misinformation, particularly in digital spaces. By fine-tuning a large language model on domain-specific datasets such as HealthFact, SciFact, and COVID-19 Fake News, the model becomes more capable of accurately evaluating medical claims. This allows it to go beyond generic language understanding and provide meaningful classifications **TRUE**, **FALSE**, or **MISLEADING** based on evidence. The ability to detect misinformation reliably can support healthcare professionals, fact-checkers, and the public in making informed decisions, especially during times of health crises.

---

## **3. API/Token Setup**

**Objective:** To authenticate securely with Hugging Face and access the Mistral-7B-Instruct-v0.1 model for fine-tuning and inference.

## Platform: Hugging Face

### 1. List the steps you followed to generate the token:

- Logged into <https://huggingface.co>
- Clicked on the profile icon → **Access Tokens**
- Clicked “**Create new token**”
- Gave it a name (colab\_token) and selected the permission scope (Read)
- Clicked “**Generate**”, then securely copied the token
- Used from huggingface\_hub import login in Colab to log in via login()
- Token was kept secure and never hardcoded

### 2. Screenshot:

**Access Tokens**

User Access Tokens + Create new token

Access tokens authenticate your identity to the Hugging Face Hub and allow applications to perform actions based on token permissions. **Do not share your Access Tokens with anyone**; we regularly check for leaked Access Tokens and remove them immediately.

Name	Value	Last Refreshed Date	Last Used Date	Permissions
colab_token	hf_...Mruc	6 days ago	4 days ago	READ

### 3. Secure Loading of Token in Code:

Used from huggingface\_hub import login in Colab to log in via login(). Token was stored only during the session and never exposed publicly

```
from huggingface_hub import login
login() # Paste your token when prompted
```

Copy a token from [your Hugging Face tokens page](#) and paste it below.  
Immediately click login after copying your token or it might be stored in plain text in this notebook file.

Token:

Add token as git credential?

Login

**Pro Tip:** If you don't already have one, you can create a dedicated 'notebooks' token with 'write' access, that you can then easily reuse for all notebooks.

#### 4. Environment Setup

To fine-tune and evaluate the Mistral-7B model, the entire workflow was implemented in a cloud-based Jupyter environment using **Google Colab Pro**. This setup allowed access to high-memory GPUs (such as A100), and supported 4-bit quantized training using bitsandbytes. All libraries were installed within the Colab environment using pip, and training states were optionally offloaded to Google Drive for efficiency.

- **Development Platform:** Google Colab Pro
- **GPU Available:** Yes
- **GPU Type:** NVIDIA A100
- **Python Version:** 3.11.12
- **Other Tools Used:**
  - Google Drive (for dataset storage and model offloading)
  - Jupyter Notebook (Colab interface)

---

#### 5. LLM Setup

For this project, the Mistral-7B-Instruct-v0.1 model was used as the base language model. It was accessed from the Hugging Face Model Hub and deployed in Google Colab with 4-bit quantization to ensure memory-efficient training. The model was wrapped with the LoRA (Low-Rank Adaptation) configuration to support parameter-efficient fine-tuning. This setup enabled us to adapt the model to domain-specific misinformation classification tasks using custom datasets while significantly reducing training cost.

I ensured compatibility and reliability by carefully managing library versions and setting up the environment in Colab with appropriate GPU support (NVIDIA A100).

- **Model Name:** Mistral-7B-Instruct-v0.1
- **Model Provider:** Hugging Face
- **Fine-Tuning Technique:** LoRA (Low-Rank Adaptation)
- **Precision Mode:** 4-bit quantization (nf4) using BitsAndBytes
- **Training Location:** Google Colab Pro with NVIDIA A100 GPU
- **Tokenizer:** AutoTokenizer from transformers
- **Model Loader:** AutoModelForCausalLM with BitsAndBytesConfig (4-bit quantization)
- **Fine-Tuning Method:** LoRA (Low-Rank Adaptation) using the peft library

## Key Python Libraries & Versions:

Library	Purpose / Description
transformers	Used to load and manage the Mistral-7B model and tokenizer. Supports model inference and generation workflows.
datasets	Used to load benchmark datasets (HealthFact, SciFact, COVID-19), format them, and apply splits for training/testing.
peft	Provides support for Parameter-Efficient Fine-Tuning (PEFT) techniques, such as LoRA, to reduce training cost.
bitsandbytes	Enables 4-bit and 8-bit quantized model loading for memory-efficient training and inference on limited GPUs.
torch (PyTorch)	Core deep learning library used to manage tensors, models, and perform training operations with GPU acceleration.
sentence-transformers	Used to compute semantic embeddings for both claims and model outputs to calculate trust scores using cosine similarity.
scikit-learn	Applied to evaluate classification performance through metrics like accuracy, precision, recall, and confusion matrix.
pandas	Used for loading and processing datasets in tabular form, including filtering, splitting, and formatting.
matplotlib	Used to visualize evaluation results such as confusion matrices and metric comparison bar charts.
gradio	Built a user-friendly web interface for testing the fine-tuned model with real-time claim predictions and trust scores.
tqdm	Provides real-time progress bars during long-running loops, such as benchmark evaluation across multiple prompts.
os	Used to manage environment variables and directories securely during setup and model loading.
warnings	Suppressed non-critical warnings during evaluation and visualization to improve notebook readability.

## Model Setup Highlights:

For this project, I worked with the **Mistral-7B-Instruct-v0.1** model, a powerful open-weight large language model suitable for instruction-based tasks. I configured the model using **4-bit quantization** (via bitsandbytes) to ensure it could be loaded efficiently in Google Colab's memory-constrained environment.

I used the Hugging Face transformers library to load the model and tokenizer, and wrapped the model with **LoRA adapters** using the peft library to support **parameter-efficient fine-tuning**. This setup enabled training with significantly fewer parameters while maintaining high performance on classification tasks.

The model was mapped to the GPU automatically using device\_map="auto", and offloading was used where necessary to balance compute across available hardware.

## Key configurations included:

- Model: mistralai/Mistral-7B-Instruct-v0.1
- Tokenizer: Loaded from the same model repository
- Quantization: 4-bit (NF4) via BitsAndBytesConfig
- PEFT technique: LoRA (Low-Rank Adaptation) with custom hyperparameters
- Fine-tuning Framework: Hugging Face Trainer

This setup ensured low-memory consumption, modular training, and ease of deployment, making it suitable for multi-stage domain-specific fine-tuning.

## Code: Install & Import

```
!pip install -q transformers datasets peft accelerate bitsandbytes scikit-learn pandas matplotlib tqdm sacrebleu rouge-score
!pip install --upgrade fsspec datasets
!pip install gradio --quiet

Preparing metadata (setup.py) ... done
      51.8/51.8 kB 3.5 MB/s eta 0:00:00
      76.1/76.1 MB 25.1 MB/s eta 0:00:00
      104.1/104.1 kB 6.6 MB/s eta 0:00:00
      61.1/61.1 kB 5.1 MB/s eta 0:00:00
      31.3/31.3 MB 72.2 MB/s eta 0:00:00
      363.4/363.4 kB 3.2 MB/s eta 0:00:00
      13.8/13.8 MB 108.7 MB/s eta 0:00:00
      24.6/24.6 MB 91.5 MB/s eta 0:00:00
      883.7/883.7 kB 57.8 MB/s eta 0:00:00
      664.8/664.8 kB 1.7 MB/s eta 0:00:00
      211.5/211.5 kB 11.8 MB/s eta 0:00:00
      56.3/56.3 kB 44.6 MB/s eta 0:00:00
      127.9/127.9 kB 9.9 MB/s eta 0:00:00
      207.5/207.5 kB 12.1 MB/s eta 0:00:00
      21.1/21.1 kB 108.5 MB/s eta 0:00:00

# STEP 2: Import Libraries
import pandas as pd
import numpy as np
from datasets import load_dataset, Dataset, concatenate_datasets
from transformers import AutoTokenizer, AutoModelForCausalLM, TrainingArguments, Trainer, DataCollatorForSeq2Seq, BitsAndBytesConfig
from peft import get_peft_model, LoraConfig, TaskType
import torch
from torch.utils.data import DataLoader
import matplotlib.pyplot as plt
from sentence_transformers import SentenceTransformer, util
import faiss
from sklearn.metrics import classification_report, confusion_matrix, ConfusionMatrixDisplay, precision_score, recall_score, f1_score
from tqdm import tqdm
from google.colab import drive
import google.colab.data_table as data_table
from sacrebleu.metrics import BLEU
from rouge_score import rouge_scorer
drive.mount('/content/drive')
import json
import os
from bert_score import score as bert_score
os.environ["WANDB_DISABLED"] = "true"
import gradio as gr

[ ] #suppress the warnings
import logging
logging.getLogger("transformers.modeling_utils").setLevel(logging.ERROR)
```

---

## 6. Dataset Description

This project utilized three key datasets, each containing medical claims along with ground truth labels to support misinformation classification tasks

### Dataset Name & Source:

- **Health Fact**
  - **Source:** Kempelen Institute of Intelligent Technologies (KInIT)
  - **URL:** <https://github.com/kinit-sk/medical-misinformation-dataset>
  - This dataset consists of annotated medical claims from reliable and unreliable sources, designed specifically for health misinformation detection.
- **SciFact**
  - **Source:** Allen Institute for AI (via Hugging Face Datasets)
  - **URL:** <https://huggingface.co/datasets/allenai/scifact>
  - It contains expert-written scientific claims along with abstracts and evidence, labeled as "SUPPORTED," "REFUTED," or "NEEDS MORE INFO." For this project, these were mapped to "TRUE," "FALSE," and "MISLEADING" for consistency.
- **COVID-19 Fake News**
  - **Source:** Kaggle
  - **URL:** <https://www.kaggle.com/datasets/arashnic/covid19-fake-news>
  - This dataset focuses on claims and tweets related to COVID-19 and classifies them as either true or false. It was further extended with manually created **MISLEADING** samples for this project to balance the label distribution.

### Feature Dictionary / Variable Description:

- **HealthFact**
  - claim\_id: Unique identifier for the claim
  - claim: The medical statement to be evaluated
  - explanation: Supporting or refuting information related to the claim
  - label: Classification label – TRUE / FALSE / MISLEADING
- **SciFact**
  - claim: Scientific medical claim
  - evidence\_text: Abstracts used to support or refute the claim
  - label: Originally SUPPORTED / REFUTED / NEI (mapped to TRUE / FALSE / MISLEADING)

- **COVID-19 Fake News**

- id: Tweet or entry identifier
- tweet: A public COVID-19 related claim
- label: TRUE / FALSE (additional MISLEADING samples created manually)

**Was preprocessing done? If yes, describe:**

Yes, comprehensive preprocessing was performed to prepare the datasets for fine-tuning and evaluation. The steps included:

- **Label normalization:** All labels were converted to lowercase and stripped of whitespace to ensure uniformity across datasets ("TRUE", "FALSE", "MISLEADING").
- **Data cleaning:** Entries with missing claims, explanations, or labels were removed to prevent inconsistencies during training.
- **Format alignment:** Datasets were converted into a unified **instruction-response format** using the ChatML structure to match the Mistral-7B input style.
- **Manual augmentation:** For the COVID-19 dataset, which originally had only TRUE and FALSE labels, 30 new **MISLEADING** samples were manually crafted and added to create a balanced three-class dataset.
- **Tokenization:** Preprocessed text was tokenized using the Mistral-compatible tokenizer with truncation and padding (max length 512), generating input IDs, attention masks, and labels.
- **Balanced splitting:** Data was split into training and testing sets using a stratified sampling strategy to preserve label distribution across both splits.

This preprocessing ensured that the model received structured, high-quality input during fine-tuning and evaluation stages.

## Code: Load & Preprocess Dataset

### Load Healthfact dataset

#### HealthFact Dataset Preprocessing

- Loaded the HealthFact data file from Google Drive.
- Selected the fields claim, explanation, and label.
- Dropped any rows with missing data.
- Normalized label values (converted to lowercase and stripped whitespace).
- Filtered the dataset to retain only the three relevant classes: "true", "false", and "misleading".
- [link text](#)Printed the final cleaned dataset size and class distribution.

```
❶ # Load HealthFact dataset
hf_df = pd.read_json("/content/drive/MyDrive/HealthFacts/healthFact_data.json", lines=True)
hf_df.head()
hf_df = hf_df[["claim", "explanation", "label"]].dropna()
# Clean and normalize
hf_df["label"] = hf_df["label"].astype(str).str.lower().str.strip()
hf_df = hf_df[hf_df["label"].isin(["true", "false", "misleading"])]
print(f"Cleaned HealthFact dataset: {len(hf_df)} entries")

hf_df["label"].value_counts()

❷ Cleaned HealthFact dataset: 1412 entries
   count
   label
   misleading    779
      false       433
      true        200
dtype: int64
```

## Load Scifi dataset

### SciFact Dataset Preprocessing

- Loaded the SciFact dataset from the Hugging Face Hub using the "claims" configuration, which contains scientific claims labeled with evidence.
- Converted the Hugging Face dataset to a pandas DataFrame for easier manipulation.
- Selected columns claim (renamed to text) and evidence\_label (renamed to label)
- Normalized the label values:"SUPPORT" to "true" and "CONTRADICT" to "false"
- Any other label dropped by assigning "unknown" and filtering out
- Printed the number of cleaned samples and the label distribution to confirm the dataset is ready for training.

```
[ ] ##from datasets import load_dataset
scifact = load_dataset("allenai/scifact", "claims", split="train")
scifact_df = scifact.to_pandas()
scifact_df = scifact_df[["claim", "evidence_label"]]
scifact_df.columns = ["text", "label"]

# Normalize labels based on actual values
def normalize_label(label):
    label = str(label).strip().lower()
    if label == "support":
        return "true"
    elif label == "contradict":
        return "false"
    else:
        return "unknown"

scifact_df["label"] = scifact_df["label"].apply(normalize_label)
scifact_df = scifact_df[scifact_df["label"] != "unknown"]

print(f"Cleaned SciFact dataset: {len(scifact_df)} entries")
print(scifact_df.head())
scifact_df["label"].value_counts()

Cleaned SciFact dataset: 957 entries
   text      label
1  1 in 5 million in UK have abnormal PrP positiv...  false
4  32% of liver transplantation programs required...  true
7  40mg/day dosage of folic acid and 2mg/day dosa...  true
8  40mg/day dosage of folic acid and 2mg/day dosa...  true
16 76-85% of people with severe mental disorder r...  true

   count
label
  true     616
  false    341

dtype: int64
```

## Merge Health Fact & SciFact Dataset

### Merging HealthFact and SciFact Datasets

This code merges the cleaned HealthFact and SciFact datasets into a unified format suitable for fine-tuning. First, the claim column in the HealthFact dataset is renamed to text to match the structure of the SciFact dataset. The explanation column is dropped since it is not used in the model training process at this stage. Both datasets are reduced to contain only the text and label columns.

The two datasets are then concatenated using pd.concat() to form a single DataFrame called combined\_df. To ensure data consistency, the code drops any rows with missing values in either the text or label fields. It then cleans the data by stripping whitespace from text, converting the text and label fields to string format, and standardizing all labels to uppercase (TRUE, FALSE, or MISLEADING).

Finally, the script prints the number of entries in the combined dataset and displays the distribution of labels. This merged and cleaned dataset becomes the foundation for the next step: formatting the data in ChatML style for supervised fine-tuning of the model.

```
[ ] # Rename for compatibility
hf_df = hf_df.rename(columns={"claim": "text"})
hf_df = hf_df[["text", "label"]]
scifact_df = scifact_df[["text", "label"]]

# Combine
combined_df = pd.concat([hf_df, scifact_df], ignore_index=True)
combined_df = combined_df.dropna(subset=["text", "label"])
combined_df["text"] = combined_df["text"].astype(str).str.strip()
combined_df["label"] = combined_df["label"].astype(str).str.strip().str.upper()

print(f"Combined dataset: {len(combined_df)} entries")
print(combined_df["label"].value_counts())
```

→ Combined dataset: 2369 entries  
label  
TRUE 816  
MISLEADING 779  
FALSE 774  
Name: count, dtype: int64

## Load Covid-19 dataset

### Load and Merge COVID-19 Dataset with MISLEADING Samples

```
❶ # Load COVID fake news dataset ---- original code
covid_ds_data = load_dataset("nancy1025/covid_fake_news", split="train")
# ---START ----sample to test---#
# Take only the first 50 examples
# covid_ds_data = full_ds.select(range(100))
# ---END----#
# Load misleading examples from Drive
misleading_path = "/content/drive/MyDrive/CovidDataset/covid_misleading.json"
df_misleading = pd.read_json(misleading_path, lines=True)
df_misleading = df_misleading.sample(50)
mis_ds = Dataset.from_pandas(df_misleading)

# Concatenate both datasets
covid_ds = concatenate_datasets([covid_ds_data, mis_ds])

#misleading test set for benchmarking and evaluation
misleading_path_test = "/content/drive/MyDrive/CovidDataset/covid_misleading_test.json"
df_misleading_test = pd.read_json(misleading_path_test, lines=True)
mis_ds_test = Dataset.from_pandas(df_misleading_test)

#Convert misleading examples into Hugging Face dataset
#mis_ds = Dataset.from_list(data_misleading)

#Concatenate both dataset
covid_ds_combined = concatenate_datasets([covid_ds_data, mis_ds])
# Take a random sample of 50 entries
covid_ds = covid_ds_combined.shuffle(seed=42)

#Verify
print(f"Random sample size: {len(covid_ds)}")
# verify data and columns
print(f"Loaded {len(covid_ds)} entries")
print(covid_ds[0])
print(covid_ds.to_pandas()["label"].value_counts())
print(full_ds_test[0])
```

→ Random sample size: 6628  
Loaded 6628 entries  
label  
real 3360  
fake 3068  
MISLEADING 290  
Name: count, dtype: int64  
{'id': 1, 'tweet': 'Our daily update is published. States reported 734k tests 39k new cases and 532 deaths. Current hospitalizations fell below 30k for the first time since June 22. <https://t.co/wzSYMe@Sh1>', 'label': 'real'}

---

## 7. Improving LLM Performance

### STAGE -1

To improve the misinformation classification capabilities of the Mistral-7B model, I implemented a structured fine-tuning pipeline comprising prompt engineering, hyperparameter tuning, dataset preparation, parameter-efficient training using LoRA, and evaluation through benchmarking and visualization.

#### Prompt Engineering & ChatML Formatting

All input examples were formatted using a function that generated ChatML-style prompts compatible with the Mistral-7B-Instruct model. Depending on availability, the function included claim explanations to provide context for classification. Prompts were wrapped in [INST]...[/INST] blocks with clearly defined classification instructions, and final labels were uppercased to maintain consistency across datasets.

```
❶ def format_chatml(row):
    claim = row.get("text", "").strip()
    explanation = row.get("explanation", "").strip()
    label = row.get("label", "").strip().upper()

    instruction = (
        "Read the following medical claim carefully. Respond only with one of these labels: TRUE, FALSE, or MISLEADING – no explanation needed.\n"
    )

    if explanation:
        prompt = (
            f"<s>[INST] {instruction}\n\nClaim: {claim}\nExplanation: {explanation} [/INST] {label} </s>"
        )
    else:
        prompt = (
            f"<s>[INST] {instruction}\n\nClaim: {claim} [/INST] {label} </s>"
        )

    return {"text": prompt}
```

#### Save Formatted Prompts to JSONL File for Fine-Tuning

```
[ ] chatml_data = combined_df.apply(format_chatml, axis=1).tolist()

with open("healthfact_balanced_3class.jsonl", "w") as f:
    for item in chatml_data:
        json.dump(item, f)
        f.write("\n")

print("ChatML JSONL saved with all 3 classes.")
```

```
✉ ChatML JSONL saved with all 3 classes.
```

#### Load Formatted ChatML Dataset for Training

```
[ ] dataset = load_dataset("json", data_files="healthfact_balanced_3class.jsonl")['train']
print(dataset)
combined_df["label"].value_counts()

Generating train split: 2369/0 [00:00<00:00, 126928.02 examples/s]
Dataset({
    features: ['text'],
    num_rows: 2369
})

  count
  label
  TRUE      816
  MISLEADING 779
  FALSE     774

dtype: int64
```

## Dataset Preparation

- **Total Samples:** 2369 entries (TRUE: 816, MISLEADING: 779, FALSE: 774)
- The data was saved as a JSONL file (healthfact\_balanced\_3class.jsonl) containing all 3 label classes.
- Dataset was split 70:30 into training and test sets, tokenized to a max length of 512 tokens.
- Labels were mapped to input\_ids for causal language modeling using Hugging Face's tokenizer.

```
[ ] tokenizer = AutoTokenizer.from_pretrained("mistralai/Mistral-7B-v0.1", use_fast=True)

# Set pad_token if missing
if tokenizer.pad_token is None:
    tokenizer.pad_token = tokenizer.eos_token

def tokenize(data_to_tokenize):
    toks = tokenizer(
        data_to_tokenize["text"],
        padding="max_length",
        truncation=True,
        max_length=512,
    )
    toks["labels"] = toks["input_ids"].copy()
    return toks

split_dataset = dataset.train_test_split(test_size=0.3, seed=42)

# Tokenize with clean output
tokenized_dataset = split_dataset.map(
    tokenize,
    batched=True,
    remove_columns=["text"]
)

tokenizer_config.json: 100% [██████████] 996/996 [00:00<00:00, 122kB/s]
tokenizer.model: 100% [██████████] 493k/493k [00:00<00:00, 18.9MB/s]
tokenizer.json: 100% [██████████] 1.80M/1.80M [00:00<00:00, 7.83MB/s]
special_tokens_map.json: 100% [██████████] 414/414 [00:00<00:00, 37.0kB/s]
Map: 100% [██████████] 1658/1658 [00:00<00:00, 3120.43 examples/s]
Map: 100% [██████████] 711/711 [00:00<00:00, 3273.62 examples/s]
```

## Model Setup and Fine-Tuning with LoRA

- **Model Used:** Mistral-7B-Instruct-v0.1
- **Precision:** 4-bit quantization (NF4) using bitsandbytes
- **LoRA Config:**
  - target\_modules: ["q\_proj", "v\_proj"]
  - r=8, lora\_alpha=16, dropout=0.05
- **Training Parameters:**
  - epochs = 3, batch\_size = 1, gradient\_accumulation\_steps = 8
  - Optimizer: AdamW with learning rate 2e-5
  - Scheduler: cosine decay
  - Checkpoints saved at each epoch
- The Trainer API from Hugging Face was used to manage training and evaluation.

```
[1]: model = AutoModelForCausalLM.from_pretrained(
    "mistralai/Mistral-7B-v0.1",
    load_in_4bit=True,
    device_map="auto"
)

lora_config = LoraConfig(
    r=8,
    lora_alpha=16,
    target_modules=["q_proj", "v_proj"],
    lora_dropout=0.05,
    bias="none",
    task_type="CAUSAL_LM"
)

model = get_peft_model(model, lora_config)
```

File	Progress	Total Size	Download Speed
config.json	100%	571/571 [00:00<00:00, 68.1kB/s]	
model.safetensors.index.json	100%	25.1k/25.1k [00:00<00:00, 492kB/s]	
Fetching 2 files	100%	2/2 [00:38<00:00, 38.19s/it]	
model-00002-of-00002.safetensors	100%	4.54G/4.54G [00:23<00:00, 160MB/s]	
model-00001-of-00002.safetensors	100%	9.94G/9.94G [00:37<00:00, 417MB/s]	
Loading checkpoint shards	100%	2/2 [00:16<00:00, 7.65s/it]	
generation_config.json	100%	116/116 [00:00<00:00, 14.2kB/s]	

```
[ ] training_args = TrainingArguments(
    output_dir=".results",
    run_name="mistral-health-scifact-stage1",
    per_device_train_batch_size=1,
    gradient_accumulation_steps=8,
    num_train_epochs=3,
    learning_rate=2e-5,
    logging_dir=".logs",
    logging_steps=10,
    save_strategy="epoch",
    bf16=True,
    fp16=not torch.cuda.is_bf16_supported(), # auto fallback if needed
    optim="adamw_torch",
    report_to="none",
    gradient_checkpointing=False, # turn OFF if GPU is spiking
    lr_scheduler_type="cosine", # smoother decay
)
```

```
[ ] # This works if you still have clean_data or split_covid

trainer = Trainer(
    model=model,
    args=training_args,
    train_dataset=tokenized_dataset["train"],
    eval_dataset=tokenized_dataset["test"],
    tokenizer=tokenizer
)

:3: FutureWarning: `tokenizer` is deprecated and will be re
  trainer = Trainer(
No label_names provided for model class `PeftModelForCausalLM`. Since `PeftModel` hides ba
```

### Start the fine-tuning process

```
[ ] trainer.train()

:3: FutureWarning: `tokenizer` is deprecated and will be re
  warnings.warn(
[621/621 24:49, Epoch 2/3]

Step Training Loss
 10      4.593500
 20      1.600100
```

### Stage 1 Results Summary

- **Training Dataset Size:** 70% of 2369 (~1658 samples)
- **Validation Dataset Size:** 30% of 2369 (~711 samples)
- **Training Loss (Final):** 0.280
- **Evaluation Loss:** 0.149
- **Epochs Completed:** 2.99
- **Eval Speed:** 15.9 samples/sec



```
eval_results = trainer.evaluate()
print("Eval Results:", eval_results)
```

[89/89 00:43] Eval Results: {'eval\_loss': 0.14919067919254303, 'eval\_runtime': 44.5004, 'eval\_samples\_per\_second': 15.977, 'eval\_steps\_per\_second': 2.0, 'epoch': 2.9891435464414}

## Evaluation on 100-Sample Benchmark Set

After completing Stage 1 fine-tuning, I evaluated the model's performance using a benchmark set of 100 samples randomly drawn from the HealthFact and SciFact datasets. This smaller sample aimed to provide a lightweight but representative snapshot of model generalization.

### Prompt Formatting for Evaluation

Each test sample was converted into a ChatML-style prompt structured as follows:

```
# Sample 100 examples from combined dataset
benchmark_df_stg1 = combined_df.sample(100, random_state=42).reset_index(drop=True)
def format_instruction(row):
    instruction = (
        "You are a helpful medical assistant. Determine if the following claim is TRUE, FALSE, or MISLEADING.\n"
        "Respond only with one of these labels: TRUE, FALSE, or MISLEADING."
    )
    return f"<s>[INST] {instruction}\n\nClaim: {row['text']} [/INST]"

# Apply to your benchmark set
eval_prompts_stg1 = benchmark_df_stg1.apply(format_instruction, axis=1).tolist()
actuals_stg1 = benchmark_df_stg1["label"].tolist() # ground truth

# Run prediction on formatted prompts
predicted_outputs_stg1 = predict_labels(model, tokenizer, eval_prompts_stg1)
```

→ Predicting: 100% [██████] 25/25 [00:25<00:00, 1.01s/it]

This instruction was designed to simulate real-world user interaction, prompting the model to return a direct label classification (TRUE, FALSE, or MISLEADING) without additional explanation.

### Batched Inference & Output Postprocessing

- I used a custom `predict_labels()` function to perform batched inference with a batch size of 4, tokenizing prompts and generating model outputs efficiently on the GPU.
- Outputs were decoded into plain text and passed through a custom `extract_label()` function, which parsed the content and extracted only the classification label (e.g., "TRUE") from inside the model-generated response.
- This ensured consistency when comparing predicted vs actual labels.

```
[ ] def predict_labels(model, tokenizer, text_list, batch_size=4):
    # Set model to evaluation mode
    model.eval()
    predictions = []

    # Iterate over the text in batches
    for i in tqdm(range(0, len(text_list), batch_size), desc="Predicting"):
        batch = text_list[i:i + batch_size]

        # Tokenize the batch
        inputs = tokenizer(
            batch,
            return_tensors="pt",
            padding=True,
            truncation=True,
            max_length=512
        )

        # Move inputs to the same device as the model
        inputs = {k: v.to(model.device) for k, v in inputs.items()}

        # Generate predictions without computing gradients
        with torch.no_grad():
            outputs = model.generate(
                input_ids=inputs["input_ids"],
                attention_mask=inputs["attention_mask"],
                pad_token_id=tokenizer.pad_token_id,
                max_new_tokens=10 # Adjust based on expected output length
            )

        # Decode the output tokens into strings
        decoded = tokenizer.batch_decode(outputs, skip_special_tokens=True)
        predictions.extend(decoded)

    return predictions
```

```
[ ] import re

def extract_label(text):

    # Step 1: Extract the part after [/INST]
    match = re.search(r"\[/INST\]\s*(.+?)\s*(</s>|\$)", text, flags=re.IGNORECASE | re.DOTALL)
    if not match:
        return "UNKNOWN"

    response = match.group(1).strip().lower()

    # Step 2: Clean markdown/special characters
    response = re.sub(r"\#[^\-:\[\]\(\)]", " ", response)
    response = re.sub(r"\s+", " ", response).strip()

    # Step 3: Match exact label words
    if re.search(r"\btrue\b", response):
        return "TRUE"
    elif re.search(r"\bfalse\b", response):
        return "FALSE"
    elif re.search(r"\bmisleading\b", response):
        return "MISLEADING"
    elif re.search(r"\bmislead\b", response):
        return "MISLEADING"
    else:
        return "UNKNOWN"
```

## Ground Truth vs Predicted Comparison

The true labels for each benchmark sample were known (from the label field), allowing a direct comparison against the extracted predictions. I computed evaluation metrics including precision, recall, and F1-score for each class using sklearn.metrics.

**Classification Metrics (on benchmark set of 100 samples):**



evaluate\_predictions(predicted\_outputs\_stg1, actuals\_stg1)



Classification Report:

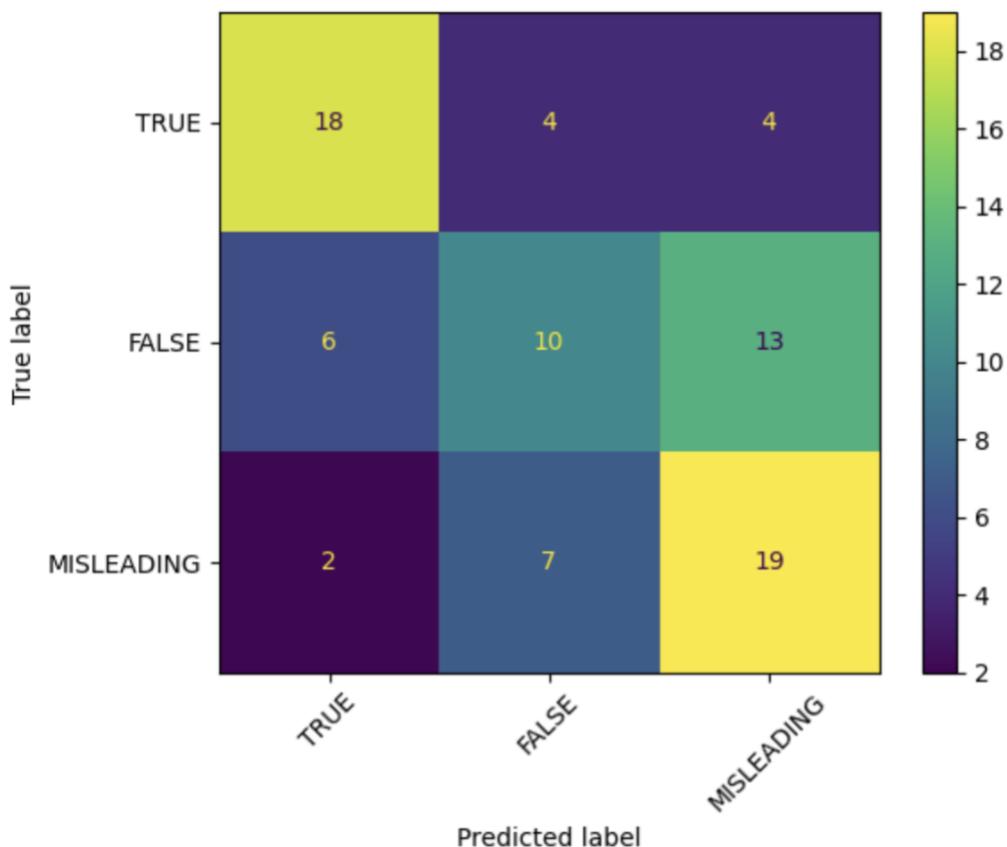
	precision	recall	f1-score	support
TRUE	0.69	0.64	0.67	28
FALSE	0.48	0.32	0.38	31
MISLEADING	0.53	0.46	0.49	41
micro avg	0.57	0.47	0.51	100
macro avg	0.57	0.48	0.51	100
weighted avg	0.56	0.47	0.51	100

## Key Observations

The benchmark results from this 100-sample evaluation indicate a reasonably balanced classification capability across all three labels. While TRUE claims showed strong precision and recall, the model was also able to identify MISLEADING and FALSE claims with moderate accuracy. All three labels recorded meaningful F1-scores, suggesting the model can generalize well when provided with clearly structured prompts. These insights help validate the effectiveness of Stage 1 fine-tuning and prompt formatting strategies in real-world misinformation detection scenarios.

## Confusion Matrix Insights

The confusion matrix shows how well the model was able to predict each class—TRUE, FALSE, and MISLEADING. Most correct predictions appear along the diagonal, where we see 18 TRUE, 10 FALSE, and 19 MISLEADING correctly classified. Some incorrect predictions are seen off the diagonal. For example, the model often confused FALSE with MISLEADING (13 cases) and MISLEADING with FALSE (7 cases). This may be because these types of claims sometimes use similar language, making them harder to separate. Overall, the model performed best with TRUE claims and showed room for improvement in handling FALSE and MISLEADING ones. This analysis helps identify where the model struggles and what can be improved in the future.



## Saving Stage 1 Evaluation Results

Following the benchmark evaluation using 100 claims from the Stage 1 dataset, I organized the results into a structured DataFrame for consistent documentation and analysis. Each row includes the original claim, its ground truth label, the model's raw output, and the final predicted label extracted via a custom parsing function. This output was exported to a CSV file (`stage1_predictions.csv`) for further comparison and storage. The dataset was then reloaded to verify integrity and later used in combined stage-wise evaluations and visualizations.

The screenshot shows a Jupyter Notebook cell with the following Python code:

```
results_df_stage1 = pd.DataFrame({
    "Claim": benchmark_df_stg1["text"], # use 'text' for Stage 1
    "Ground Truth": actuals_stg1,
    "Model Output": predicted_outputs_stg1,
    "Predicted Label": [extract_label(p) for p in predicted_outputs_stg1]
})
results_df_stage1.to_csv("stage1_predictions.csv", index=False)
results_df_stage1_data = pd.read_csv("stage1_predictions.csv")
data_table.DataTable(results_df_stage1_data)
```

Below the code is a screenshot of a Jupyter Notebook displaying a DataFrame titled "stage1\_predictions.csv". The DataFrame has four columns: "index", "Claim", "Ground Truth", "Model Output", and "Predicted Label". There are 100 rows of data. The "Model Output" column contains the raw JSON responses from the model, and the "Predicted Label" column contains the extracted binary labels (TRUE, FALSE, or MISLEADING).

index	Claim	Ground Truth	Model Output	Predicted Label
0	New valve procedure doesn't open heart	FALSE	[INST] You are a helpful medical assistant. Determine if the following claim is TRUE, FALSE, or MISLEADING. Respond only with one of these labels: TRUE, FALSE, or MISLEADING. Claim: New valve procedure doesn't open heart [/INST] FALSE [REASON] The procedure is called	FALSE
1	Planned Parenthood opts out of U.S. subsidies in fight over abortion referrals.	TRUE	[INST] You are a helpful medical assistant. Determine if the following claim is TRUE, FALSE, or MISLEADING. Respond only with one of these labels: TRUE, FALSE, or MISLEADING. Claim: Planned Parenthood opts out of U.S. subsidies in fight over abortion referrals. [/INST] TRUE ➔	TRUE
2	In rhesus macaques, daily subcutaneous injections of tenofovir protects against rectally transmitted simian-human immunodeficiency virus.	TRUE	[INST] You are a helpful medical assistant. Determine if the following claim is TRUE, FALSE, or MISLEADING. Respond only with one of these labels: TRUE, FALSE, or MISLEADING. Claim: In rhesus macaques, daily subcutaneous injections of tenofovir protects against rectally transmitted simian-human immunodeficiency virus. [/INST] TRUE ➔	TRUE

## STAGE 2: Fine-Tuning and Evaluation on COVID-19 Dataset

In Stage 2, I extended the fine-tuning process by incorporating a domain-specific dataset comprising COVID-19 fake news samples along with manually crafted synthetic MISLEADING entries. This stage aimed to enhance the model's accuracy in distinguishing subtle forms of misinformation within pandemic-related claims.

### Dataset Preparation

For Stage 2 fine-tuning, I curated a more domain-specific dataset focused on pandemic-related misinformation. The goal was to improve the model's ability to detect nuanced and misleading COVID-19 claims. This phase involved combining **COVID-19 Fake News samples** with a **set of synthetic misleading claims**, ensuring coverage across all three target labels: TRUE, FALSE, and MISLEADING.

### Data Sources:

- **COVID-19 Fake News Dataset:** This public dataset includes various real-world tweets related to COVID-19. Each entry contains a tweet and a label indicating whether the claim is true or false.
- **Synthetic MISLEADING Samples:** Since the original dataset lacked sufficient misleading examples, I created additional examples based on known misinformation patterns and public health myths. These were designed to simulate ambiguous or partially correct claims, which are typically harder to detect.

```

# Load COVID fake news dataset ---- original code
covid_ds_data = load_dataset("nanny1025/covid_fake_news", split="train")
full_ds_test = load_dataset("nanny1025/covid_fake_news", split="test")
# --START ----sample to test---#
# Take only the first 50 examples
#covid_ds_data = full_ds.select(range(100))
#-----END-----#
# Load misleading examples from Drive
misleading_path = "/content/drive/MyDrive/CovidDataset/covid_misleading.json"
df_misleading = pd.read_json(misleading_path, lines=True)
# Convert to Hugging Face Dataset
mis_ds = Dataset.from_pandas(df_misleading)
# Concatenate both datasets
covid_ds = concatenate_datasets([covid_ds_data, mis_ds])

#Misleading test set for benchmarking and evaluation
misleading_path_test = "/content/drive/MyDrive/CovidDataset/covid_misleading_test.json"
df_misleading_test = pd.read_json(misleading_path_test, lines=True)
mis_ds_test = Dataset.from_pandas(df_misleading_test)

#Convert misleading examples into Hugging Face dataset
#mis_ds = Dataset.from_list(data_misleading)

#Concatenate both dataset
covid_ds_combined = concatenate_datasets([covid_ds_data, mis_ds])
# Take a random sample of 50 entries
covid_ds = covid_ds_combined.shuffle(seed=42)

# Verify
print(f"Random sample size: {len(covid_ds)}")
# verify data and columns
print(f"Loaded {len(covid_ds)} entries")
print(covid_ds[0])
print(covid_ds.to_pandas()["label"].value_counts())
print(full_ds_test[0])

```

```

Random sample size: 6620
Loaded 6620 entries
{'id': 931, 'tweet': 'A video of an overcrowded Shramik train claimed to be from Mumbai travelling to West Bengal.', 'label': 'fake', 'explanation': None}
label
real      3360
fake      3060
MISLEADING    200
Name: count, dtype: int64
{'id': 1, 'tweet': 'Our daily update is published. States reported 734k tests 39k new cases and 532 deaths. Current hospitalizations fell below 30k for the first ti

```

#### Inspect COVID-19 Dataset in Pandas

```

[ ] # bring your dataset into pandas
pdf = covid_ds.to_pandas()

# show the first few labels and their unique set
print("Sample labels:", pdf["label"].head(10).tolist())
print("All unique labels:", pdf["label"].unique())

```

```

Sample labels: ['fake', 'real', 'real', 'fake', 'fake', 'fake', 'real', 'MISLEADING', 'fake']
All unique labels: ['fake' 'real' 'MISLEADING']

```

## Prompt Engineering & ChatML Formatting

Each tweet was transformed into a structured instruction format using a custom function. Prompts followed the ChatML convention, wrapping claims within [INST]...[/INST] blocks. The prompt instructed the model to classify the input claim as either **TRUE**, **FALSE**, or **MISLEADING**, enforcing strict one-word responses to reduce ambiguity.

### Format COVID Dataset into ChatML Structure for Fine-Tuning

This code formats the normalized COVID-19 dataset into the ChatML prompt-response format expected by instruction-tuned models like Mistral. Each entry is saved in a .jsonl file, making it ready for tokenization and training.

```
[ ] def format_chatml_binary(example):
    instruction = (
        "You are a helpful medical assistant. Determine if the following claim is TRUE, FALSE, or MISLEADING.\n"
        "Respond only with one of these labels: TRUE, FALSE, or MISLEADING."
    )

    claim = example.get("tweet", "").strip()
    label = example.get("label", "").strip().upper()

    return {
        "text": (
            f"<s>[INST] {instruction}\n\nClaim: {claim} [/INST] {label} </s>"
        )
    }

# Apply and dump
chatml_binary = covid_ds.map(format_chatml_binary)
with open("covid19_binary_chatml.jsonl", "w") as out:
    for rec in chatml_binary:
        json.dump(rec, out)
        out.write("\n")

print("Saved covid19_binary_chatml.jsonl with TRUE/FALSE/MISLEADING labels.")
```

### Load and Clean COVID ChatML Dataset for Tokenization

This block loads the previously saved covid19\_binary\_chatml.jsonl into a Hugging Face dataset and strips away any unused metadata columns to retain only the formatted text. The result is a clean dataset ready for tokenization and fine-tuning.

```
[ ] covid_data = load_dataset("json", data_files="covid19_binary_chatml.jsonl")["train"]
print(covid_data.column_names)
# Drop all original columns, keep only the newly formatted "text"
clean_data = covid_data.map(
    format_chatml_binary,
    remove_columns=covid_data.column_names # this rips out id, tweet, label, etc.
)
print(clean_data.column_names) # should output: ['text']

→ Generating train split: 6620/0 [00:00<00:00, 258951.12 examples/s]
['id', 'tweet', 'label', 'explanation', 'text']
Map: 100% 6620/6620 [00:00<00:00, 15397.89 examples/s]
['text']
```

## Tokenization for Stage 2 Fine-Tuning

ChatML-style prompt were tokenized using AutoTokenizer with a maximum length of 512 tokens, producing input\_ids and attention\_mask. The tokenized sequences were structured to ensure only the label portion was used for loss computation by masking the rest of the input with -100. This tokenization setup enabled efficient autoregressive learning during LoRA-based fine-tuning, ensuring that the model learned to predict only the classification label while preserving context from the input claim

```

def tokenize_for_clm(covidData):
    toks = tokenizer(
        covidData["text"],
        padding="max_length",
        truncation=True,
        max_length=512,
    )
    # key trick: for causal LM we train the model to reproduce the entire input,
    # so the labels are just a copy of input_ids
    toks["labels"] = toks["input_ids"].copy()
    return toks

# Tokenize with clean output

covid_tokenized = clean_data.map(
    tokenize_for_clm,
    batched=True,
    remove_columns=["text"], # now only input_ids, attention_mask, labels remain
)

print(covid_tokenized.column_names)
# should be ['input_ids','attention_mask','labels']
print("Before tokenization:", len(clean_data)) # Should match jsonl
print("After tokenization:", len(covid_tokenized))
!head -n 3 covid19_binary_chatml.jsonl

```

Map: 100% 6620/6620 [00:01<00:00, 3675.54 examples/s]

['input\_ids', 'attention\_mask', 'labels']

Before tokenization: 6620

After tokenization: 6620

{"id": 931, "tweet": "A video of an overcrowded Shramik train claimed to be from Mumbai travelling to West Bengal.", "label": "FALSE", "explanation": null, "text": null}, {"id": 1171, "tweet": "\u201cAs we mark 100 days since our first case of #COVID19Nigeria we remember over 300 people we\u2019ve lost to this outbreak. As we publish", "label": "TRUE", "explanation": null, "text": null}, {"id": 182, "tweet": "New COVIDView reports hospitalization rates and deaths from #COVID19 are increasing. Indicators that track hospitalizations and deaths usually"},

## Stage 2 Fine-Tuning with LoRA

For Stage 2, I extended the fine-tuning of the Mistral-7B-Instruct-v0.1 model using the peft library with LoRA (Low-Rank Adaptation). The model was re-wrapped with get\_peft\_model() and configured for causal language modeling. Training was conducted using the Hugging Face Trainer API over three epochs with a learning rate of 2e-4, batch size of 2, and gradient accumulation steps set to 4. Mixed precision (fp16) was used to optimize GPU usage. The training output was stored in a dedicated directory, with logs and checkpoints maintained per epoch. This setup improved training efficiency while preserving earlier knowledge from Stage 1.

- lora\_dropout: Prevents overfitting within the LoRA path.

```

from peft import LoraConfig, get_peft_model, TaskType

model = get_peft_model(model, lora_config) # Re-wrap with LoRA
model.config.use_cache = False

# /usr/local/lib/python3.11/dist-packages/peft/mapping_func.py:73: UserWarning: You are trying to modify a model with PEFT for a second time. If you want to reload it
#     warnings.warn(
# /usr/local/lib/python3.11/dist-packages/peft/mapping_func.py:79: UserWarning: The PEFT config's 'base_model_name_or_path' was renamed from 'mistralai/Mistral-7B-v0'.
#     warnings.warn(
# /usr/local/lib/python3.11/dist-packages/peft/tuners/tuners_utils.py:167: UserWarning: Already found a 'peft_config' attribute in the model. This will lead to having
#     warnings.warn(

```



```

    training_args_stage2 = TrainingArguments([
        output_dir='./mistral_stage2_covid_binary',
        per_device_train_batch_size=2,
        gradient_accumulation_steps=4,
        num_train_epochs=3,
        learning_rate=2e-4,
        logging_dir='./logs',
        logging_steps=10,
        save_strategy='epoch',
        fp16=True,
        save_total_limit=2,
        report_to='none',
        remove_unused_columns=False
    ])
    #split_covid = covid_tokenized.train_test_split(test_size=0.3, seed=42)
    trainer_stg2 = Trainer(
        model=model, # This is existing LoRA-wrapped Stage 1 model
        args=training_args_stage2,
        train_dataset=covid_tokenized,
        eval_dataset=covid_tokenized["test"],
        tokenizer=tokenizer,
    )
    trainer_stg2.train()

    >>> <ipython-input-40-a1402e1d0de9>:16: FutureWarning: 'tokenizer' is deprecated and will be removed in version 5.0.0 for `Trainer.__init__`. Use `processing_class` instead.
    No label_names provided for model class `PeftModelForCausalLM`. Since `PeftModel` hides base models input arguments, if label_names is not given, label_names can't be used.
    [2481/2481 1:10.08, Epoch 2/3]
    Step  Training Loss
    10      1.733200
    ..      1.733200

```

```

] model.save_pretrained("mistral_stage2_covid_binary")
tokenizer.save_pretrained("mistral_stage2_covid_binary")

-> ('mistral_stage2_covid_binary/tokenizer_config.json',
     'mistral_stage2_covid_binary/special_tokens_map.json',
     'mistral_stage2_covid_binary/tokenizer.model',
     'mistral_stage2_covid_binary/added_tokens.json',
     'mistral_stage2_covid_binary/tokenizer.json')

```

## Stage 2 Benchmarking and Label Extraction

To evaluate Stage 2 model performance on COVID-19 misinformation, I created a benchmark set of 100 examples by combining 80 tweets from the COVID-19 Fake News dataset with 20 synthetic MISLEADING samples. Each claim was formatted using ChatML-style prompts with a fixed instruction. The formatted prompts were tokenized and passed to the fine-tuned Mistral-7B model.

Model responses were post-processed to extract the predicted label using a multi-step parsing method. This involved stripping out the assistant's reply after the instruction block ([/INST]), removing special characters, normalizing whitespace, and matching the content against expected labels ("TRUE", "FALSE", or "MISLEADING"). A fallback logic was also used to catch near matches or ambiguous cases. This ensured robust label extraction across varied model responses and enabled accurate performance reporting. The predictions were later compared against ground-truth labels to generate evaluation metrics.

```

❸ covid_ds_test_combined = concatenate_datasets([full_ds_test.select(range(100)), mis_ds_test])
benchmark_df_covid = covid_ds_test_combined.to_pandas().sample(100, random_state=42).reset_index(drop=True)
print(covid_ds_test_combined[0])
def format_prompt(row):
    instruction = (
        "You are a helpful medical assistant. Determine if the following claim is TRUE, FALSE, or MISLEADING.\n"
        "Respond only with one of these labels: TRUE, FALSE, or MISLEADING."
    )
    claim = row.get("tweet", "").strip()
    return f"<s>[INST] {instruction}\n\nClaim: {claim} [/INST]"
eval_texts = benchmark_df_covid.apply(format_prompt, axis=1).tolist()
actuals = benchmark_df_covid["label"].tolist()

❹ {'id': 1, 'tweet': 'Our daily update is published. States reported 734k tests 39k new cases and 532 deaths. Current hospitalizations fell below 30k for the first t

```

### Label Extraction & Standardization

```

[ ] def extract_label_from_response(text):
    text = text.strip()

    # Step 1: Extract text after [/INST]
    match_inst = re.search(r"\[/INST\]\s*(.*?)\s*(</s>|\$)", text, flags=re.IGNORECASE | re.DOTALL)
    if match_inst:
        response = match_inst.group(1).strip().lower()

    # Step 2: Remove markdown/special chars
    response = re.sub(r"[*#>\-:\[\]\(\)]", " ", response)
    response = re.sub(r"\s+", " ", response).strip()

    # Step 3: Match exact labels
    if re.fullmatch(r"true", response):
        return "TRUE"
    elif re.fullmatch(r"false", response):
        return "FALSE"
    elif re.fullmatch(r"misleading", response):
        return "MISLEADING"

    # Step 4: Fuzzy fallback
    for lbl in ["true", "false", "misleading"]:
        if lbl in response:
            return lbl.upper()

    # Step 5: Final fallback – search anywhere
    text = text.lower()
    for lbl in ["true", "false", "misleading"]:
        if lbl in text:
            return lbl.upper()

    return "UNKNOWN"

```

## Benchmarking Evaluation – Stage 2

To assess the fine-tuned model's performance in handling targeted COVID-19 misinformation, a benchmark evaluation was conducted using a curated 100-sample dataset combining real and synthetic examples. Each input was structured using ChatML prompts, and predicted responses were post-processed to extract the final labels for evaluation.

## Classification Performance:

```

print(classification_report(actuals, predicted_labels_stg2))

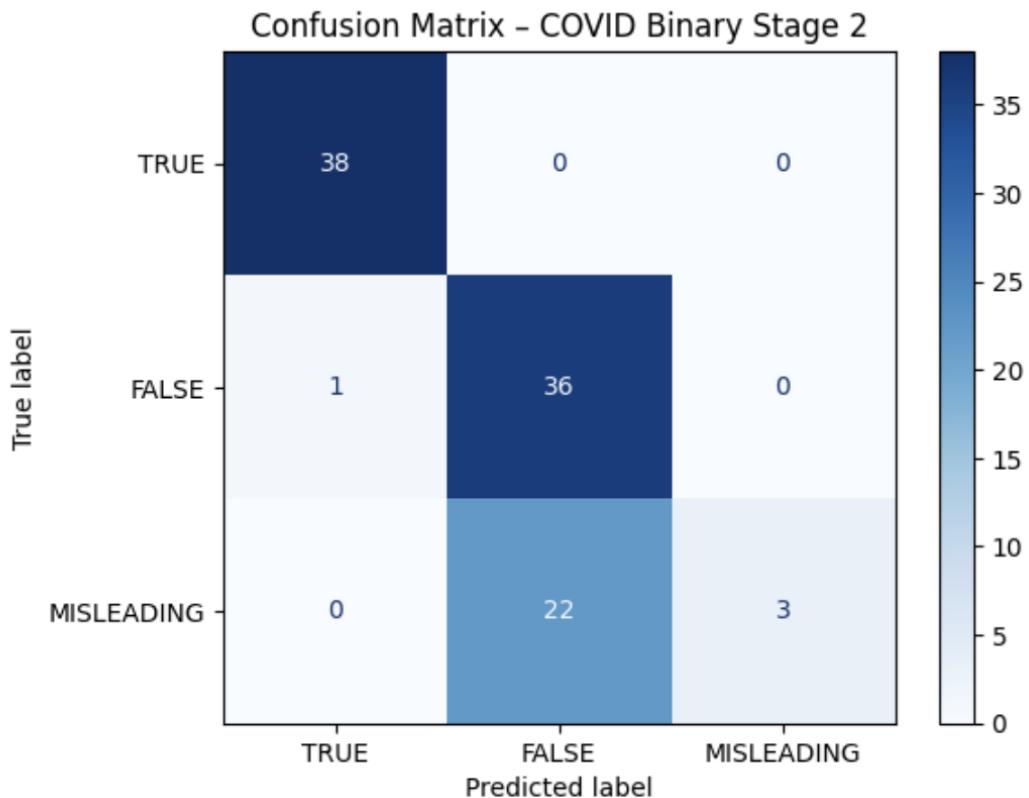
cm = confusion_matrix(actuals, predicted_labels_stg2, labels=["TRUE", "FALSE", "MISLEADING"])
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=["TRUE", "FALSE", "MISLEADING"])
disp.plot(cmap="Blues")
plt.title("Confusion Matrix - COVID Binary Stage 2")
plt.show()

```

	precision	recall	f1-score	support
FALSE	0.62	0.97	0.76	37
MISLEADING	1.00	0.12	0.21	25
TRUE	0.97	1.00	0.99	38
accuracy			0.77	100
macro avg	0.87	0.70	0.65	100
weighted avg	0.85	0.77	0.71	100

### Key Observations:

- The model showed strong performance on TRUE and FALSE claims.
- MISLEADING label recall remained low, indicating difficulty in identifying subtle or ambiguous misinformation.
- Overall accuracy reached **77%**, with weighted F1-score at **0.71**, suggesting meaningful improvements post fine-tuning over the specialized COVID dataset.



The confusion matrix presents a visual summary of how the Stage 2 fine-tuned model performed on a benchmark set of 100 COVID-related claims. Out of 38 claims labeled as TRUE, the model correctly predicted all 38. Similarly, for FALSE claims, 36 out of 37 were classified accurately. However, the model struggled with MISLEADING claims—only 3 out of 25 were predicted correctly, while 22 were misclassified as FALSE.

This result shows that while the model was highly confident in detecting TRUE and FALSE statements, it had difficulty recognizing MISLEADING claims. This may be due to the subtle or overlapping nature of misleading content, which often lacks clear falsehoods. The confusion matrix confirms the need for further improvements, such as increasing the number of diverse misleading samples during training or applying techniques like RAG to support better evidence-based reasoning.

## Saving Stage 2 Evaluation Results

After completing the evaluation of the Stage 2 fine-tuned Mistral-7B model, I created a structured DataFrame to log all evaluation outputs for reproducibility and further analysis. This included:

- **Claim:** The original COVID-19 tweet used as input.
- **Ground Truth:** The correct label as per the benchmark dataset (TRUE, FALSE, or MISLEADING).
- **Model Output:** The raw text response generated by the model.
- **Predicted Label:** The final predicted classification extracted using a rule-based label extraction function.

These results were saved as a CSV file (covid\_stage2\_predictions.csv) and previewed using a pandas DataTable. This export ensures transparency in reporting and enables traceable validation of model behavior on real-world misinformation samples.

```
[ ] results_df_covid = pd.DataFrame({
    "Claim": benchmark_df_covid["tweet"],
    "Ground Truth": actuals,
    "Model Output": predicted_outputs,
    "Predicted Label": predicted_labels_stg2
})
```

### Saving the Evaluation Results

```
(results_df_covid.to_csv("covid_stage2_predictions.csv", index=False)
print("Saved results to covid_stage2_predictions.csv")
results_df_covid_data = pd.read_csv("covid_stage2_predictions.csv")
data_table.DataTable(results_df_covid_data)
```

⤒ Saved results to covid\_stage2\_predictions.csv

index	Claim	Ground Truth	Model Output	1 to 25 of 100 entries	Filter	?	Predicted Label
0	Tablik Jamaat Chief Maulana Saad: "If 70000 people get coronavirus India will be destroyed. The country will be in our possession."	FALSE	[INST] You are a helpful medical assistant. Determine if the following claim is TRUE, FALSE, or MISLEADING. Respond only with one of these labels: TRUE, FALSE, or MISLEADING. Claim: Tablik Jamaat Chief Maulana Saad: "If 70000 people get coronavirus India will be destroyed. The country will be in our possession." [INST] FALSE				FALSE
1	Two interesting correlations: 1) Children tend to weather COVID-19 pretty well; they also get a ton of Vitamin D. 2) Black people are getting slammed by COVID-19; black people also have much higher instances of Vitamin D deficiency (76% vs 40% in the general population).	FALSE	[INST] You are a helpful medical assistant. Determine if the following claim is TRUE, FALSE, or MISLEADING. Respond only with one of these labels: TRUE, FALSE, or MISLEADING. Claim: Two interesting correlations: 1) Children tend to weather COVID-19 pretty well; they also get a ton of Vitamin D. 2) Black people are getting slammed by COVID-19; black people also have much higher instances of Vitamin D deficiency (76% vs 40% in the general population). [INST] FALSE				FALSE
2	This is the sixth time a global health emergency has been declared under the International Health Regulations but it is easily the most severe-@DrTedoros https://t.co/JvKCOPtett	TRUE	[INST] You are a helpful medical assistant. Determine if the following claim is TRUE, FALSE, or MISLEADING. Respond only with one of these labels: TRUE, FALSE, or MISLEADING. Claim: This is the sixth time a global health emergency has been declared under the International Health Regulations but it is easily the most severe-@DrTedoros https://t.co/JvKCOPtett [INST] TRUE				TRUE
3	Eating garlic prevents coronavirus infection.	MISLEADING	[INST] You are a helpful medical assistant. Determine if the following claim is TRUE, FALSE, or MISLEADING. Respond only with one of these labels: TRUE, FALSE, or MISLEADING. Claim: Eating garlic prevents coronavirus infection. [INST] FALSE				FALSE

## Trust Score Computation Using Sentence Embeddings

To complement the model's classification performance with a measure of semantic confidence, I computed a **Trust Score** for each prediction using the Sentence-BERT model all-MiniLM-L6-v2. This metric evaluates how closely the model's generated response aligns with the original input claim.

The process involved:

- **Embedding Generation:** I used the SentenceTransformer to generate vector embeddings for each claim and its corresponding model output.
- **Cosine Similarity:** For each pair, I calculated the cosine similarity between the embeddings. This provides a score between -1 and 1, where higher values indicate stronger semantic alignment.
- **Score Extraction:** Only the diagonal elements (each claim with its own response) were extracted to form the final Trust Score.
- **Integration:** These scores were appended to the evaluation DataFrame for interpretability and analysis.

```
[ ] # Load a small but effective embedding model
embedder = SentenceTransformer("all-MiniLM-L6-v2")
```

→ dules.json: 100% [██████████] 349/349 [00:00<00:00, 42.4kB/s]

fig\_sentence\_transformers.json: 100% [██████████] 116/116 [00:00<00:00, 14.6kB/s]

ADME.md: 100% [██████████] 10.5k/10.5k [00:00<00:00, 1.26MB/s]

entence\_bert\_config.json: 100% [██████████] 53.0/53.0 [00:00<00:00, 6.08kB/s]

fig.json: 100% [██████████] 612/612 [00:00<00:00, 77.3kB/s]

: Storage is enabled for this repo, but the 'hf\_xet' package is not installed. Falling back to regular HTTP download. For bet  
NING:huggingface\_hub.file\_download:Xet Storage is enabled for this repo, but the 'hf\_xet' package is not installed. Falling

del.safetensors: 100% [██████████] 90.9M/90.9M [00:00<00:00, 231MB/s]

enizer\_config.json: 100% [██████████] 350/350 [00:00<00:00, 45.0kB/s]

ab.txt: 100% [██████████] 232k/232k [00:00<00:00, 15.7MB/s]

enizer.json: 100% [██████████] 466k/466k [00:00<00:00, 2.02MB/s]

cial\_tokens\_map.json: 100% [██████████] 112/112 [00:00<00:00, 13.6kB/s]

fig.json: 100% [██████████] 190/190 [00:00<00:00, 24.9kB/s]

```

▶ # Use input Claim and Model Output columns
claims = results_df_covid["Claim"].tolist()
responses = results_df_covid["Model Output"].tolist()

# Compute embeddings in batches
claim_embeddings = embedder.encode(claims, convert_to_tensor=True)
response_embeddings = embedder.encode(responses, convert_to_tensor=True)

# Compute cosine similarity
cosine_scores = util.cos_sim(claim_embeddings, response_embeddings)

# Convert diagonal of cosine similarity matrix into list
trust_scores = [round(score.item(), 3) for score in cosine_scores.diag()]

# Add to DataFrame
results_df_covid["Trust Score"] = trust_scores

# Preview
results_df_covid[["Claim", "Predicted Label", "Trust Score"]].head()

```

→

	Claim	Predicted Label	Trust Score
0	Tablik Jamaat Chief Maulana Saad: "If 70000 pe...	FALSE	0.594
1	Two interesting correlations:\n\n1) Children t...	FALSE	0.651
2	This is the sixth time a global health emergen...	TRUE	0.578
3	Eating garlic prevents coronavirus infection.	FALSE	0.416
4	The coronavirus is only as dangerous as a flu ...	FALSE	0.561

### Evaluation Results with Trust Scores

```
[ ] results_df_covid.to_csv("covid_stage2_with_trust_score.csv", index=False)
print(" Saved with trust scores to covid_stage2_with_trust_score.csv")
```

---

## Final Evaluation

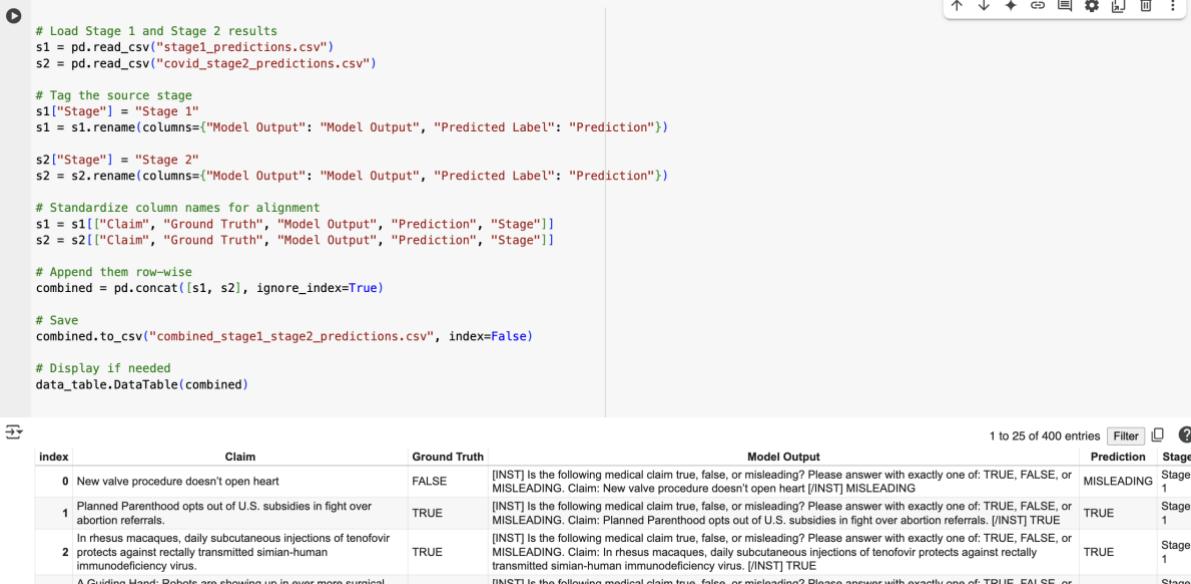
### Merging Stage 1 and Stage 2 Results

To support holistic performance analysis, I combined prediction outputs from both Stage 1 and Stage 2 into a single consolidated file. This allowed me to compare model behavior across general medical claims and COVID-19 misinformation scenarios.

Key steps:

- **Loading Data:** Imported the prediction outputs saved as stage1\_predictions.csv and covid\_stage2\_predictions.csv.

- **Tagging Stages:** Added a Stage column to each dataset to distinguish between general medical (Stage 1) and COVID-specific (Stage 2) benchmarks.
- **Standardization:** Renamed columns for consistency (Prediction, Model Output, etc.) and aligned them to a common schema: ["Claim", "Ground Truth", "Model Output", "Prediction", "Stage"].
- **Concatenation:** Used pandas.concat() to merge both datasets into a unified DataFrame.
- **Exporting Combined Results:** Saved the final merged file as combined\_stage1\_stage2\_predictions.csv.



```
# Load Stage 1 and Stage 2 results
s1 = pd.read_csv("stage1_predictions.csv")
s2 = pd.read_csv("covid_stage2_predictions.csv")

# Tag the source stage
s1["Stage"] = "Stage 1"
s1 = s1.rename(columns={"Model Output": "Model Output", "Predicted Label": "Prediction"})

s2["Stage"] = "Stage 2"
s2 = s2.rename(columns={"Model Output": "Model Output", "Predicted Label": "Prediction"})

# Standardize column names for alignment
s1 = s1[["Claim", "Ground Truth", "Model Output", "Prediction", "Stage"]]
s2 = s2[["Claim", "Ground Truth", "Model Output", "Prediction", "Stage"]]

# Append them row-wise
combined = pd.concat([s1, s2], ignore_index=True)

# Save
combined.to_csv("combined_stage1_stage2_predictions.csv", index=False)

# Display if needed
data_table.DataTable(combined)
```

1 to 25 of 400 entries Filter ?

index	Claim	Ground Truth	Model Output	Prediction	Stage
0	New valve procedure doesn't open heart	FALSE	[INST] Is the following medical claim true, false, or misleading? Please answer with exactly one of: TRUE, FALSE, or MISLEADING. Claim: New valve procedure doesn't open heart [/INST] MISLEADING	MISLEADING	Stage 1
1	Planned Parenthood opts out of U.S. subsidies in fight over abortion referrals.	TRUE	[INST] Is the following medical claim true, false, or misleading? Please answer with exactly one of: TRUE, FALSE, or MISLEADING. Claim: Planned Parenthood opts out of U.S. subsidies in fight over abortion referrals. [/INST] TRUE	TRUE	Stage 1
2	In rhesus macaques, daily subcutaneous injections of tenofovir protect against rectally transmitted simian-human immunodeficiency virus.	TRUE	[INST] Is the following medical claim true, false, or misleading? Please answer with exactly one of: TRUE, FALSE, or MISLEADING. Claim: In rhesus macaques, daily subcutaneous injections of tenofovir protects against rectally transmitted simian-human immunodeficiency virus. [/INST] TRUE	TRUE	Stage 1
	A Guirlande Hand: Robots are shown in in ever more surreal		[INST] Is the following medical claim true, false, or misleading? Please answer with exactly one of: TRUE, FALSE, or MISLEADING. Claim: A Guirlande Hand: Robots are shown in in ever more surreal		Stage 1

## Final Benchmarking

To assess the comparative performance of the fine-tuned models across both stages, I loaded the combined benchmark predictions and evaluated them separately for Stage 1 and Stage 2. This step provided clarity on how the model performed across general and domain-specific misinformation tasks.

### Steps performed:

- **Data Loading:** Imported the merged file combined\_stage1\_stage2\_predictions.csv containing predictions and labels from both stages.
- **Label Normalization:** Converted all label fields (Ground Truth and Prediction) to uppercase for consistency during evaluation.
- **Dataset Segregation:** Split the combined dataset into two subsets: Stage 1 (general medical claims) and Stage 2 (COVID-19 misinformation).
- **Classification Reports:** Used sklearn.metrics.classification\_report() to compute precision, recall, F1-score, and support for each label category across the two stages.

This dual evaluation enabled a clearer understanding of strengths and limitations in both contexts and helped justify the effectiveness of multi-stage fine-tuning. The final comparison also supports decisions on future dataset augmentation and model refinement strategies.

```
# Load and Fix Labels

# Load combined Stage 1 + Stage 2 data
df = pd.read_csv("combined_stage1_stage2_predictions.csv")

# Normalize to uppercase
df["Ground Truth"] = df["Ground Truth"].astype(str).str.upper()
df["Prediction"] = df["Prediction"].astype(str).str.upper()

# Optional: Check label consistency
print(df["Stage"].value_counts())
print(df[["Ground Truth", "Prediction", "Stage"]].dtypes)

# Separate Stage 1 and Stage 2
df_s1 = df[df["Stage"] == "Stage 1"]
df_s2 = df[df["Stage"] == "Stage 2"]

# Run Classification Reports
report_s1 = classification_report(
    df_s1["Ground Truth"],
    df_s1["Prediction"],
    labels=["TRUE", "FALSE", "MISLEADING"],
    output_dict=True
)

report_s2 = classification_report(
    df_s2["Ground Truth"],
    df_s2["Prediction"],
    labels=["TRUE", "FALSE", "MISLEADING"],
    output_dict=True
)

# Display summary
print("Stage 1 Summary")
print(classification_report(df_s1["Ground Truth"], df_s1["Prediction"], labels=["TRUE", "FALSE", "MISLEADING"]))

print("Stage 2 Summary")
print(classification_report(df_s2["Ground Truth"], df_s2["Prediction"], labels=["TRUE", "FALSE", "MISLEADING"]))
```

## Final Evaluation and Comparison of Stage 1 and Stage 2

To evaluate the effectiveness of the staged fine-tuning approach, I benchmarked both Stage 1 and Stage 2 models using 100 labeled samples each. The results were analyzed using standard classification metrics—precision, recall, and F1-score—across the three output classes: TRUE, FALSE, and MISLEADING.

### Stage 1 (HealthFact + SciFact)

The Stage 1 model demonstrated moderate ability to classify TRUE claims (F1-score: 0.67), with lower performance on FALSE (F1: 0.38) and MISLEADING (F1: 0.49). These results suggest the model was able to identify factual claims better than deceptive or ambiguous ones. The weighted average F1-score was 0.51, and the model showed a micro average accuracy of 57%.

### Stage 2 (COVID-19 + Synthetic MISLEADING)

Fine-tuning on domain-specific data led to notable performance improvements. The Stage 2 model reached an F1-score of 0.99 for TRUE and 0.76 for FALSE, indicating enhanced understanding of verifiable and inaccurate claims. However, despite perfect precision, the MISLEADING class had a recall of only 0.12 (F1: 0.21), highlighting ongoing challenges in

detecting nuanced misinformation. The overall weighted average F1-score improved to 0.71, with 77% accuracy on the benchmark set.

## Conclusion

The comparison confirms that Stage 2 fine-tuning substantially improved classification accuracy for TRUE and FALSE labels. While MISLEADING detection remains difficult, the overall performance gains affirm the value of domain-adaptive instruction fine-tuning in health misinformation tasks.

---

→ Stage

Stage 1	100			
Stage 2	100			
Name:	count, dtype: int64			
Ground Truth	object			
Prediction	object			
Stage	object			
dtype:	object			
Stage 1 Summary				
	precision	recall	f1-score	support
TRUE	0.69	0.64	0.67	28
FALSE	0.48	0.32	0.38	31
MISLEADING	0.53	0.46	0.49	41
micro avg	0.57	0.47	0.51	100
macro avg	0.57	0.48	0.51	100
weighted avg	0.56	0.47	0.51	100
Stage 2 Summary				
	precision	recall	f1-score	support
TRUE	0.97	1.00	0.99	38
FALSE	0.62	0.97	0.76	37
MISLEADING	1.00	0.12	0.21	25
accuracy			0.77	100
macro avg	0.87	0.70	0.65	100
weighted avg	0.85	0.77	0.71	100

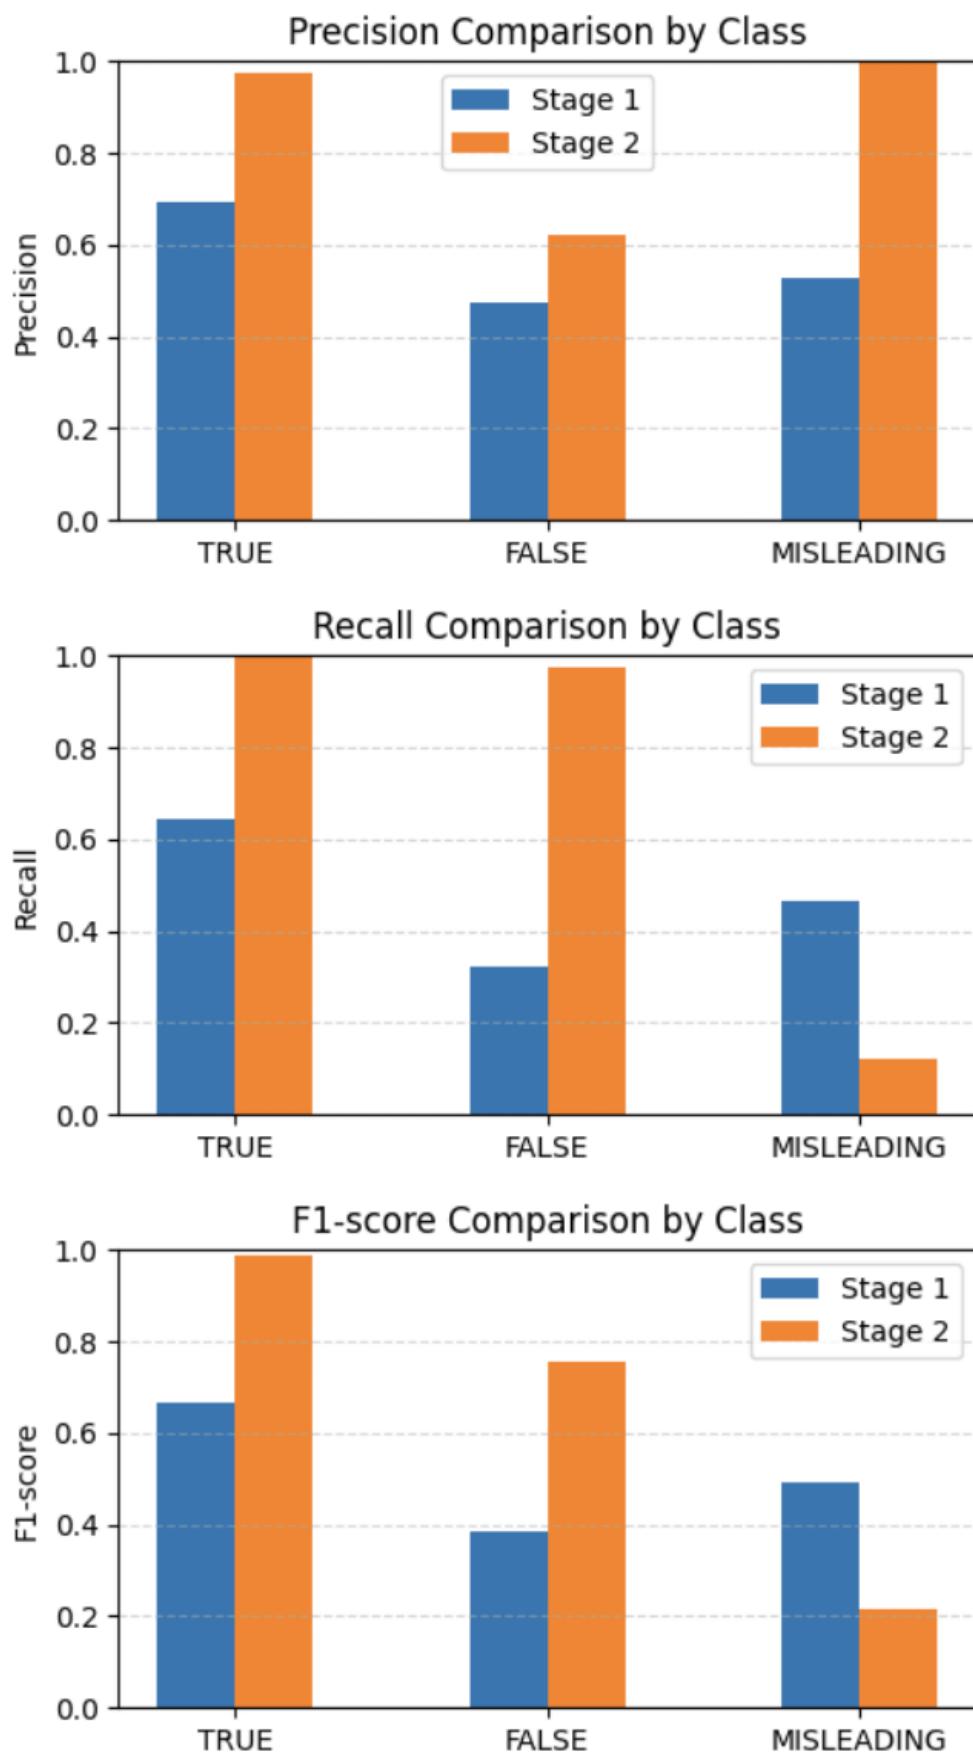
## Insights from Evaluation Metrics

The updated precision, recall, and F1-score comparison charts highlight the improvements gained through Stage 2 fine-tuning. The precision for the FALSE class increased from 0.48 to 0.62, while the TRUE class rose from 0.69 to 0.97—indicating significantly stronger confidence in predictions post fine-tuning. Although the MISLEADING class reached perfect precision (1.00), its recall dropped from 0.46 to 0.12. This means the model confidently identifies misleading claims when it predicts them, but still fails to detect many such instances.

Recall for the FALSE class improved from 0.32 to 0.97, showing the effectiveness of domain-specific fine-tuning on COVID-19 misinformation. The macro and weighted averages also show a consistent improvement across all three metrics, with weighted F1-score rising from 0.51 to 0.71. These findings validate the effectiveness of a staged fine-tuning pipeline using LoRA, where Stage 2's domain-adaptive training resulted in better generalization and stronger classification, particularly for TRUE and FALSE claims.

### Figure: Precision, Recall, and F1-Score Comparison for Stage 1 vs Stage 2

This bar chart illustrates the comparative classification performance of the Mistral-7B model after Stage 1 and Stage 2 fine-tuning. Stage 2 shows a noticeable improvement in precision and recall for the TRUE and FALSE classes, confirming enhanced model reliability in identifying accurate and inaccurate claims. However, the model continues to struggle with MISLEADING claims, where recall dropped despite high precision. This visual supports the effectiveness of domain-adaptive tuning while also highlighting ongoing challenges in multi-class health misinformation classification.



---

## 8. UI Integration

To make the model accessible and interactive, a Gradio-based web interface was developed to classify COVID-19 medical claims and visualize the model's prediction confidence.

### Tools Used

- **Gradio:** An open-source Python library used to create customizable, web-based machine learning demos. It allows seamless connection between models and users through an intuitive front-end interface. Gradio was used to build the input/output form, label visualizations, and real-time result display.
- **Hugging Face Transformers:** Provided the foundation for running the fine-tuned Mistral-7B-Instruct model, allowing efficient handling of tokenization, inference, and output generation.
- **SentenceTransformers (all-MiniLM-L6-v2):** Used to compute semantic embeddings of user-entered claims and the model's output to generate a **Trust Score** based on cosine similarity.
- **Google Colab:** Served as the cloud-based development and execution environment. The UI was hosted and launched within Colab notebooks, which provided access to GPU resources and allowed sharing through a temporary Gradio public URL.

### Key Features

- **Text Input Support:** Users can input free-form COVID-19 claims (e.g., “COVID-19 can be cured by drinking lime juice”).
- **Three-Way Classification:** The interface classifies each claim as either **TRUE**, **FALSE**, or **MISLEADING** using a model fine-tuned on domain-specific datasets.
- **Visual Feedback:** The prediction result is displayed clearly with a labeled output box, along with an explanation of the decision logic in plain language.
- **Trust Score Output:** A numerical trust score (ranging from 0 to 1) is calculated and shown, offering a basic indication of how semantically aligned the model's response is to the input claim.
- **Real-Time Performance:** The Gradio app responds quickly to user queries and can be deployed with a shareable public link for testing by peers, mentors, or stakeholders.
- **Accessibility:** The app requires no installation and can be accessed directly from any browser, making it easy to test the model without coding knowledge.

```

# Load trust scorer
embedder = SentenceTransformer("all-MiniLM-L6-v2")

# Prompt formatter
def format_input_prompt(claim):
    instruction = (
        "You are a helpful medical assistant. Determine if the following claim is TRUE, FALSE, or MISLEADING.\n"
        "Respond only with one of these labels: TRUE, FALSE, or MISLEADING."
    )

    return f"<s>[INST] {instruction}\n\nClaim: {claim} [/INST]"

# UI enabling function
def classify_claim(claim):
    # Format the prompt for the model
    prompt = format_input_prompt(claim)

    # Run inference
    output = predict_labels(model, tokenizer, [prompt])[0]

    # Extract label
    label = extract_label_from_response(output)

    # Compute trust score
    claim_embedding = embedder.encode(claim, convert_to_tensor=True)
    response_embedding = embedder.encode(output, convert_to_tensor=True)
    trust_score = util.cos_sim(claim_embedding, response_embedding).item()

    # Format a cleaner display output
    display_text = f"Your claim:\n\n{claim}\nis classified as → **{label}**"

    return label, display_text, round(trust_score, 3)

```

→ The following layers were not sharded: embeddings.token\_type\_embeddings.weight, embeddings.LayerNorm.weight, embeddings.LayerN

```

[ ] interface = gr.Interface(
    fn=classify_claim,
    inputs=gr.Textbox(lines=4, label="Enter a COVID-19 Claim"),
    outputs=[gr.Label(label="Predicted Label"),
             gr.Markdown(label="Explanation"),
             gr.Number(label="Trust Score (0-1)")]
    ),
    title="COVID-19 Misinformation Classifier (Mistral-7B)",
    description="This model classifies COVID-related claims as TRUE, FALSE, or MISLEADING using a fine-tuned Mistral-7B model.",
)

interface.launch()

```

## How It Works:

The UI wraps the fine-tuned Mistral-7B model behind a Gradio interface. When a user submits a claim, it is formatted using a custom prompt template designed for instruction-following LLMs. The model processes this prompt and returns a classification. Simultaneously, a trust score is computed using cosine similarity between the input claim and the model's generated output using all-MiniLM-L6-v2 embeddings. This score provides a rough measure of response confidence. The interface then displays the predicted label, explanation text, and trust score in a structured layout that is easy for both technical and non-technical users to understand.

## Screen Shots:

Colab notebook detected. To show errors in colab notebook, set debug=True in launch()  
! output actions ing on public URL: <https://08d1da8b237002b3a3.gradio.live>

This share link expires in 1 week. For free permanent hosting and GPU upgrades, run `gradio deploy`

This model classifies COVID-related claims as TRUE, FALSE, or MISLEADING using a fine-tuned Mistral-7B model.

Enter a COVID-19 Claim

Insulin effects appetite via ventral tegmental neurons.

Predicted Label

**TRUE**

Your claim:  
"Insulin effects appetite via ventral tegmental neurons."  
is classified as → **TRUE**

Trust Score (0–1)

0.424

**Flag**

Use via API 🚀 · Built with Gradio 🎨 · Settings ⚙️

It looks like you are running Gradio on a hosted a Jupyter notebook. For the Gradio app to work, sha  
→ Colab notebook detected. To show errors in colab notebook, set debug=True in launch()  
\* Running on public URL: <https://08d1da8b237002b3a3.gradio.live>

This share link expires in 1 week. For free permanent hosting and GPU upgrades, run `gradio deploy`

## COVID-19 Misinformation Classifier (Mistral-7B)

This model classifies COVID-related claims as TRUE, FALSE, or MISLEADING using a fine-tuned Mistral-7B model.

Enter a COVID-19 Claim

Taking a hot bath will prevent you from catching COVID-19

Predicted Label

**FALSE**

Your claim:  
"Taking a hot bath will prevent you from catching COVID-19"  
is classified as → **FALSE**

Trust Score (0–1)

0.522

**Flag**

→ It looks like you are running Gradio on a hosted a Jupyter notebook. For the Gradio app to work, sha

Colab notebook detected. To show errors in colab notebook, set debug=True in launch()  
\* Running on public URL: <https://08d1da8b237002b3a3.gradio.live>

This share link expires in 1 week. For free permanent hosting and GPU upgrades, run `gradio deploy`

This model classifies COVID-related claims as TRUE, FALSE, or MISLEADING using a fine-tuned Mistral-7B model.

Enter a COVID-19 Claim

COVID-19 can be cured by chanting or spiritual rituals.

Predicted Label

**MISLEADING**

Your claim:

"COVID-19 can be cured by chanting or spiritual rituals."

is classified as → **MISLEADING**

Trust Score (0–1)

0.444

Flag

Use via API 🚀 · Built with Gradio 🎨 · Settings 🌐

→ It looks like you are running Gradio on a hosted a Jupyter notebook. For the Gradio app to work, sha

Colab notebook detected. To show errors in colab notebook, set debug=True in launch()  
\* Running on public URL: <https://08d1da8b237002b3a3.gradio.live>

This share link expires in 1 week. For free permanent hosting and GPU upgrades, run `gradio deploy`

This model classifies COVID-related claims as TRUE, FALSE, or MISLEADING using a fine-tuned Mistral-7B model.

Enter a COVID-19 Claim

COVID-19 can be cured by drinking lime juice.

Predicted Label

**MISLEADING**

Your claim:

"COVID-19 can be cured by drinking lime juice."

is classified as → **MISLEADING**

Trust Score (0–1)

0.429

Flag

Use via API 🚀 · Built with Gradio 🎨 · Settings 🌐